

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/309224333>

# A containerized mesoscale model and analysis toolkit to accelerate classroom learning, collaborative research, and uncertainty...

Article in *Bulletin of the American Meteorological Society* · October 2016

DOI: 10.1175/BAMS-D-15-00255.1

---

CITATIONS

0

READS

13

8 authors, including:



[Joshua P. Hacker](#)

National Center for Atmospheric Research

77 PUBLICATIONS 816 CITATIONS

[SEE PROFILE](#)



[Ivo Jimenez](#)

University of California, Santa Cruz

16 PUBLICATIONS 8 CITATIONS

[SEE PROFILE](#)



[Carlos Maltzahn](#)

University of California, Santa Cruz

86 PUBLICATIONS 1,391 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



The MATERHORN Project [View project](#)



# AMERICAN METEOROLOGICAL SOCIETY

*Bulletin of the American Meteorological Society*

## EARLY ONLINE RELEASE

This is a preliminary PDF of the author-produced manuscript that has been peer-reviewed and accepted for publication. Since it is being posted so soon after acceptance, it has not yet been copyedited, formatted, or processed by AMS Publications. This preliminary version of the manuscript may be downloaded, distributed, and cited, but please be aware that there will be visual differences and possibly some content differences between this version and the final published version.

The DOI for this manuscript is doi: 10.1175/BAMS-D-15-00255.1

The final published version of this manuscript will replace the preliminary version at the above DOI once it is available.

If you would like to cite this EOR in a separate work, please use the following full citation:

Hacker, J., J. Exby, D. Gill, I. Jimenez, C. Maltzahn, T. See, G. Mullendore, and K. Fossell, 2016: A containerized mesoscale model and analysis toolkit to accelerate classroom learning, collaborative research, and uncertainty quantification. Bull. Amer. Meteor. Soc. doi:10.1175/BAMS-D-15-00255.1, in press.



# A containerized mesoscale model and analysis toolkit to accelerate

## classroom learning, collaborative research, and uncertainty quantification

Joshua P. Hacker\* John Exby, David Gill

*Research Applications Laboratory, National Center for Atmospheric Research, Boulder, CO.*

Ivo Jimenez, Carlos Maltzahn

*Computer Science Department, University of California Santa Cruz*

Timothy See, Gretchen Mullendore

*Department of Atmospheric Sciences, University of North Dakota*

Kathryn Fossell

*Mesoscale Microscale Meteorology Division, National Center for Atmospheric Research*

\*Corresponding author address: Joshua P. Hacker, Research Applications Laboratory, NCAR, P.O.

Box 3000, Boulder, CO 80307.

E-mail: hacker@ucar.edu

## ABSTRACT

14 Numerical weather prediction (NWP) experiments can be complex and time  
15 consuming; results depend on computational environments and numerous in-  
16 put parameters. Delays in learning and obtaining research results are in-  
17 evitable. Students face disproportionate effort in the classroom or begin-  
18 ning graduate-level NWP research. Published NWP research is generally  
19 not reproducible, introducing uncertainty and slowing efforts that build on  
20 past results. This work exploits the rapid emergence of software container  
21 technology to produce a transformative research and education environment.

22 The Weather Research and Forecasting (WRF) model anchors a set of linked  
23 Linux-based containers, which include software to initialize and run the  
24 model, analyze results, and serve output to collaborators. The containers are  
25 demonstrated with a WRF simulation of Hurricane Sandy. The demonstration  
26 illustrates the following: (1) how the often-difficult exercise in compiling the  
27 WRF and its many dependencies is eliminated ; (2) how sharing containers  
28 provides identical environments for conducting research; (3) that numerically  
29 reproducible results are easily obtainable; and (4) how uncertainty in the re-  
30 sults can be isolated from uncertainty arising from computing system differ-  
31 ences. Numerical experiments to simultaneously measure numerical repro-  
32 ducibility and sensitivity to compiler optimization provide guidance for inter-  
33 preting NWP research. Reproducibility is independent from operating system  
34 and hardware. Results here show numerically identical output on all comput-  
35 ing platforms tested. Performance reproducibility is also demonstrated. The  
36 result is an infrastructure to accelerate classroom learning, graduate research,  
37 and collaborative science.

<sup>38</sup> **Capsule Summary**

<sup>39</sup> Software containers can revolutionize research and education with numerical weather prediction  
<sup>40</sup> models by easing use and guaranteeing reproducibility.

<sup>41</sup> **1. Introduction**

<sup>42</sup> Numerical models are a cornerstone of weather prediction today, and also support a broad range  
<sup>43</sup> of weather research. Studies quantifying the ability of numerical models to predict atmospheric  
<sup>44</sup> state, and simulate atmospheric phenomena, form two key lines of inquiry. By establishing model  
<sup>45</sup> fidelity to the atmosphere, those studies also provide a basis for broader scientific inquiry with  
<sup>46</sup> models. The use of numerical models in atmospheric research has become ubiquitous over the past  
<sup>47</sup> few decades. Areas of research include physical process identification and analysis, atmospheric  
<sup>48</sup> predictability, and predictions of future climates, among others.

<sup>49</sup> Extensive use of numerical models in research demands educating students for diligent appli-  
<sup>50</sup> cation of these large and complex codes. Beyond basic theory and best practices in numerical  
<sup>51</sup> methods, the opportunities for mistakes are vast. While mistakes can provide positive learning  
<sup>52</sup> experiences, they can also produce misleading results and consequently incorrect interpretation.

<sup>53</sup> The purpose of this article is to describe the implementation of software containers for nu-  
<sup>54</sup> merical weather prediction (NWP) research and education. Container technology has profound  
<sup>55</sup> implications for education and research in numerical weather prediction. Containers not only en-  
<sup>56</sup> able reproducibility, they greatly lower barriers for accessing cutting edge NWP codes. The tools  
<sup>57</sup> discussed herein, such as source code repositories and containers, may be relatively new for many  
<sup>58</sup> readers. These tools are presented in detail to fully describe the procedures used. But a deep  
<sup>59</sup> understanding is not needed. The containers presented here have already been tested successfully  
<sup>60</sup> by multiple users, primarily from universities. The technical details can remain mostly transparent

61 for new users; the power of these tools can then be further realized by users (including students  
62 and educators) who become more advanced.

63 The Weather Research and Forecasting (WRF) model (Skamarock et al. 2008) is a state-of-the-  
64 art numerical weather prediction (NWP) model for operations and research, with users numbering  
65 in the tens of thousands. It was engineered to be portable and easy to install on a limited yet  
66 common set of platforms, with code that is both complex and extensive (over 1.5 million lines).

67 It has many dependencies on external software packages (e.g. for I/O, parallel communications,  
68 and data compression) that are not trivial to satisfy; compilation and execution can be an intensive  
69 effort for beginning users who lack support from experienced system administrators. The WRF  
70 help-desk requests are dominated by new users reporting difficulty compiling the model. A new  
71 version has been released twice annually for the last 15 years, and users who want to update and  
72 rebuild may face repeated challenges with each new set of code.

73 Running the model can also be difficult for new users. Individual steps to generate computational  
74 grids, import initialization data, produce initial and boundary conditions, and run the model can  
75 be daunting. Those steps can be scripted into a workflow that can range from simple to complex,  
76 depending on the application. Scripted workflows have been developed by countless individuals  
77 and groups over the last decade; the proverbial “wheel” has been reinvented countless times, es-  
78 pecially when considering the number of individuals who have written analysis tools to compute  
79 the same diagnostics on output.

80 Implementation of experiential learning with numerical weather prediction codes is especially  
81 challenging for users in the classroom. Some universities are able to provide hands-on exercises  
82 due to a combination of sufficient lab resources and, most importantly, staff (faculty and support)  
83 with a sufficient knowledge base in modeling and information technology. But even those univer-

<sup>84</sup> sites with sufficient knowledge struggle with continual updates to compilers, operating systems,  
<sup>85</sup> and model code.

<sup>86</sup> Results from software that implements floating point operations is generally not numerically  
<sup>87</sup> reproducible. Examples of the impacts on numerical weather models can be found in Thomas  
<sup>88</sup> et al. (2002) and Baker et al. (2015). A consequence is that NWP research is not reproducible.

<sup>89</sup> Numerical uncertainty can lead to misleading results. The chaotic nature of the models means  
<sup>90</sup> that small perturbations, perhaps arising from truncation errors, can organize and grow into fea-  
<sup>91</sup> tures that are tempting to interpret as an important response to a change in input or model for-  
<sup>92</sup> mulation. The classic paper by Lorenz (1963) is an example of how numerical truncation can  
<sup>93</sup> lead to chaotic results, and in that case it was truncation that led to Lorenz's discovery of chaos.

<sup>94</sup> Many times the differences appearing meaningful cannot be reproduced, or simply cannot be dis-  
<sup>95</sup> tinguished from chance.

<sup>96</sup> Finally, collaboration amongst researchers working on different computing platforms can be  
<sup>97</sup> cumbersome. It often relies on careful, manual, code management to make sure collaborators are  
<sup>98</sup> working with the same code. Configuration of a complex model must also be done carefully, and  
<sup>99</sup> still the opportunity to make mistakes is always present.

<sup>100</sup> Software containers, which are becoming an essential part of modern software development and  
<sup>101</sup> deployment, offer a path for mitigating or eliminating many of the problems in NWP research and  
<sup>102</sup> education that are described above. In the remaining text, we describe how and show some exam-  
<sup>103</sup> ples. First, software containers are reviewed. Then in section 3 a set of containers for initializing,  
<sup>104</sup> running, and analyzing results from the WRF model are presented. Numerical and performance  
<sup>105</sup> reproducibility are demonstrated in section 4 with a simulation of Hurricane Sandy, leading to  
<sup>106</sup> a clear way to perform uncertainty quantification in section 5. Section 6 concludes with some

<sup>107</sup> discussion of the implications for container technology in numerical simulation and prediction.  
<sup>108</sup> Finally in section 7, links to on-line content and open repositories are provided.

<sup>109</sup> **2. Software containers**

<sup>110</sup> A container is a software-based packaging and distribution tool that collects all elements and  
<sup>111</sup> dependencies of a Linux-based application. Containers store the run-time environment together  
<sup>112</sup> with one or more applications for ease of transportation, installation and execution across a variety  
<sup>113</sup> of operating systems (Linux, Mac, Windows). Containers avoid any need for recompilation or  
<sup>114</sup> complex virtual machine environments. Because they only house the necessary components to  
<sup>115</sup> execute an application, they are much simpler and smaller than virtual machines running complete  
<sup>116</sup> operating systems. Containers allow an application (or suite of applications) to be developed,  
<sup>117</sup> tested and executed in such a way that when relocated to an alternative compute environment, the  
<sup>118</sup> execution and results of the application workflow are consistent and repeatable. Modern container  
<sup>119</sup> tools have origins based on Linux containers that have long been part of Linux kernels. They are  
<sup>120</sup> now evolving toward a standards-based “engine” that allows portability and compatibility despite  
<sup>121</sup> changes to underlying compute systems, networks, or operating systems.

<sup>122</sup> A image, or base image, is a pre-built set of application elements: binaries, libraries, config-  
<sup>123</sup> uration files, and any ancillary tools required to provide an entire execution environment for the  
<sup>124</sup> application to function. An image can be developed over time by being tagged and maintained  
<sup>125</sup> with version control methods, and downloaded from official repositories as a basis for developer  
<sup>126</sup> or operational use. An application image may be maintained by a known set of developers for the  
<sup>127</sup> purposes of official verification at known version levels, and contain official patches and security  
<sup>128</sup> updates.

129 Launching a particular instance of an application image into a run-time state produces a con-  
130 tainer running that application. The container is executing the entire environment of the appli-  
131 cation, without requiring the host operating system to have any software installed, other than the  
132 container engine itself.

133 Linux containers have been used in cloud services for over 10 years to provide stability, scal-  
134 ability, and a more secure, isolated run-time environment than is often possible on commonly used  
135 computing platforms. Today's container engines have defined an emerging standard, which can  
136 execute pre-built codes within every modern operating system (OS) including Windows, Mac OS  
137 X, and Linux. Containerized applications provide a well-documented command set that users can  
138 implement quickly without regard to a destination OS. Software deployment methods are simple,  
139 and reduce the level of introduction needed to bring new users or developers into a collaborative  
140 group.

141 Industry experts believe containers will continue to be prevalent, with accelerating adoption  
142 due to best-fit scaling on desktops and laptops. Commercial and research cloud platforms are  
143 providing greater off-the-shelf server resources as hardware prices either fall, or computing power  
144 per dollar spent rises. Container technology is a natural fit to commercial cloud services; a user  
145 can immediately deploy software without any knowledge of the underlying hardware or OS. Once  
146 software tools are in containers, users have easy access to vast resources of cloud computing,  
147 which can be exploited when conditions demand.

148 We chose Docker ([www.docker.com](http://www.docker.com)) as the container engine for this work, but others are avail-  
149 able. Docker containers are user-friendly, already widely adopted across the software industry,  
150 and allow users to execute complex pre-compiled scientific codes anywhere. The combination of  
151 all these features enables users to focus on running numerically reproducible scientific analysis,  
152 and to reduce start up times. Within containers, software configurations can be easily integrated

<sup>153</sup> into arbitrary workflows that assemble trivially portable modules. The containers may use differ-  
<sup>154</sup> ing data sets assigned at launch, and be built on demand. The resulting output can be stored with  
<sup>155</sup> metadata using local storage, proven distributed storage technologies, cloud storage such as S3  
<sup>156</sup> and Glacier offered by Amazon Web Services ([aws.amazon.com](http://aws.amazon.com)), or similar resources.

<sup>157</sup> The open-source community and the rapid adoption of Docker (and similar) container technolo-  
<sup>158</sup> gies has boomed since 2012, with entire companies such as Uber, Netflix, Goldman Sachs, and  
<sup>159</sup> others leveraging containers across hybrid (on-premise and cloud) platforms for dynamic scal-  
<sup>160</sup> ability that can be based on users (or events) worldwide. Frameworks for managing complex and  
<sup>161</sup> multiple instances of containers are evolving to provide streamlined tools that allow flexibility,  
<sup>162</sup> verification, reporting, and coordination.

<sup>163</sup> Containers internally run a very light-weight Linux environment, as opposed to a full Linux  
<sup>164</sup> kernel supporting an operating system. Prior to June 2016, the stable released Docker engine soft-  
<sup>165</sup> ware components for Windows and Mac OS X required an additional thin virtual machine (VM)  
<sup>166</sup> layer managed by VirtualBox ([www.virtualbox.org](http://www.virtualbox.org)), to pass container commands to a Linux VM  
<sup>167</sup> environment that supports containers. The VM layer added a slight performance loss (approx.  
<sup>168</sup> 5% in run time), and it created extra docker-machine configuration and management steps often  
<sup>169</sup> not familiar to general users. As of June 2016, new (v.1.10+ beta and beyond) versions of the  
<sup>170</sup> Docker engine for Windows and Mac OS X have bypassed the requirement of a Linux VM, com-  
<sup>171</sup> pletely discarding Virtualbox and all local Docker-machine management and VM configura-  
<sup>172</sup> tion steps. Docker containers now run natively on Windows and OS X platforms, eliminating the slight  
<sup>173</sup> performance penalty of VMs and delivering container commands to native Windows and OS X  
<sup>174</sup> kernels. Docker continues to build upon this new beta software for the purposes of incorporating  
<sup>175</sup> the same native features across Linux, Windows, and OS X.

176 **3. A set of containers for numerical weather prediction**

177 The Docker containers described above have been implemented so that a user can easily con-  
178 figure and complete a WRF run. An example deployment of the containers is summarized in the  
179 schematic Fig. 1. The WRF-Docker container also includes the WRF and WRF Pre-processing  
180 System (WPS) to ensure software version consistency and compatibility. A separate container with  
181 the NCAR Command Language (NCL; [www.ncl.ucar.edu](http://www.ncl.ucar.edu)) plotting scripts is also implemented.  
182 Easy run-time configuration is provided by external access to the WPS and WRF namelists on the  
183 local filesystem. Output and namelists (to record provenance) are written to the local file system,  
184 and additional metadata can be easily included. Pre-built data containers (including WPS and  
185 WRF namelists) for specific simulation periods and regions are currently available in open repos-  
186 itories. In that case the local folders for the namelist, input files, and static geographic (GEOG)  
187 data are unnecessary. Containers for acquiring input data, archiving output, and sharing output, are  
188 under development in collaboration with several universities under the NSF-funded Big Weather  
189 Web project ([bigweatherweb.org](http://bigweatherweb.org)).

190 The WPS/WRF container currently holds the WPS and WRF code, and pre-built binary execut-  
191 ables (applications). The container includes executables to create domains, populate files with ge-  
192 ographical and land-surface data, interpolate input meteorological data to the WRF computational  
193 grid, and run the WRF to produce a simulation or forecast. Gnu compilers and all dependencies  
194 are included. A user launches the container with command line options for recompiling the code  
195 and executing the various WPS and WRF executables. It is simple for users who want to make  
196 code modifications to enter the container and make changes before recompiling, or to recompile  
197 within the container before exiting and executing WRF runs from outside.

198     Figure 1 is an example of one possible workflow. It is easily scripted or run from the command  
199     line. The intent is for users to be able to deploy any number of the containers, and substitute his  
200     own at any step in the workflow. For example we can imagine that a user may want to introduce a  
201     different and more sophisticated post-processing approach that allows analysis supporting a physi-  
202     cal process study. That new container would be deployed in exactly the same way as long as it only  
203     needs access to the WRF output files. The results would be archived or shared in the same way,  
204     and the new container can be uploaded to the repository for community use and reproducibility.

205     In the spirit of a community and open-source effort, users can contribute useful containers so that  
206     others can reproduce results or address other scientific questions with similar tools. The question  
207     a researcher often faces, of developing his own code or trying to adapt and port code developed by  
208     others, will go away in many cases because no porting effort is required.

209     Source codes for the model and ancillary systems, the containers, and data sets (e.g. the ge-  
210     ography data needed to design WRF domains) are managed in version-controlled source reposi-  
211     tories, which allow continuous integration and unit testing through the chain of source commits  
212     to Docker-build image generation. We chose GitHub (GitHub; [www.github.com](http://www.github.com)), but could have  
213     equally chosen another repository host such as Bitbucket ([www.bitbucket.org](http://www.bitbucket.org)) or GitLab ([gitlab.com](http://gitlab.com)). The GitHub repositories are for developers, or researchers who want to rebuild contain-  
214     ers to address specific scientific questions. When a project on GitHub is executed to produce a  
215     Docker image, the resulting image is pushed to DockerHub, which manages Git repositories for  
216     Docker images. We could have equally chosen an Amazon or Google registry service for hosting  
217     the Docker images. A user who is not interested in developing (e.g. changing the WRF code or  
218     analysis tools) can simply pull images and use them to run containers with any necessary run-time  
219     input files (namelists, gridded data, etc). Keeping all known fixes and build environments updated,  
220     patched and tested through a quality assurance process via a core development team, including

222 code commits to a GitHub repository and pulls from a DockerHub image repository, guarantees  
223 software integrity in the launch environment for model results to be replicated, or queried, for  
224 future studies.

225 Collaboration on research, which may or may not require code changes, is easily enabled via  
226 Github, DockerHub, and the reproducibility properties inherent to working in containers. Collab-  
227 orators can maintain a repository with version control on the code and container. Information and  
228 links on how to obtain the source for the containers, or the images, on open GitHub and Dock-  
229 erHub repositories, is given in Section 7. The version control system ensures bit-perfect results  
230 can be obtained across native OS platforms, also that users are accessing an identical code base as  
231 required. All collaborators in a group can be sure they are using identical codes and even binaries,  
232 providing unambiguous interpretation of quantitative analysis of output from the model or set of  
233 analysis tools.

#### 234 **4. Numerical and performance reproducibility**

235 One of the primary reasons that NWP research has historically been numerically non-  
236 reproducible is that different computer architectures and environments perform mathematical op-  
237 erations differently. Differences can be greater when parallel operations are included. First, differ-  
238 ent chips have different binary truncation, leading to different effects of round-off error. Second,  
239 compilers sometimes reorder operations in attempts to speed computation; the reordering depends  
240 on the compiler brand, version, and level of optimization. Third, unless carefully implemented,  
241 parallel operations may not be performed in the order intended. Often this is not a requirement,  
242 but truncation guarantees that some computations are not reproducible when the order of a parallel  
243 computation is not enforced. For example, the result of an average depends on the precision in the  
244 sum and division, and in which order they are performed. [Thomas et al. \(2002\)](#) demonstrates the

245 effect of compiler optimization and parallel topology on NWP, and Baker et al. (2015) shows how  
246 lack of expected numerical uncertainty can indicate deficient code quality. Experiments described  
247 here test uncertainty introduced by different hardware and operating systems, parallel topology,  
248 and compiler options.

249 A computationally inexpensive WRF simulation of Hurricane Sandy is the basis for experimen-  
250 tation. The simulation is on a single computational  $50 \times 50$  horizontal domain with  $\Delta X = 40$  km  
251 and 60 vertical levels. Initial and boundary conditions are interpolated from NCEP’s Global Fore-  
252 cast System (GFS). Twelve-hour simulations are initialized at 1200 UTC 27 Oct 2012. Details of  
253 the complete sub-grid and forcing schemes (known as “physics”) are unimportant here, but they do  
254 represent a reasonable set with components widely found in the literature. This computationally  
255 inexpensive simulation serves to illustrate the capability. A plot of total precipitation accumulated  
256 over the 12-h simulation, shown in Fig. 2, is also useful to see the computational domain. The very  
257 same container demonstrated here can be trivially used to execute other WRF simulations simply  
258 by changing the input data and namelists.

259 The Sandy container simulation has been tested on dozens of different computers with a variety  
260 of hardware and operating systems. Table 1 shows a small representative sample, including com-  
261 mon Mac and Linux platforms. It also includes resources available on Amazon’s EC2 and Packet’s  
262 ([www.packet.net](http://www.packet.net)) Tiny Atom cloud computing infrastructure. Successful deployment and testing  
263 is far more extensive than shown in Table 1.

264 In each case, the same container is deployed and executed. Hourly WRF output files are saved  
265 on a user’s file system by externally linking to the WRF output directory within the container’s  
266 file system. An arbitrary platform is chosen as a reference, which serves as comparison for all of  
267 the other test deployments. A deployment in OS X 10.10.5 (Yosemite), with an Intel Xeon E5 on  
268 a Mac Pro, serves as the reference here. Free gnu compilers, namely gfortran and gcc, are part of

269 the container for building the WRF executable. Because the compiler and compiler flags are the  
270 same, the binary WRF executable is identical regardless of what computer compiles it. But it is  
271 not necessarily obvious that the same binary will produce the same output on different hardware.

272 Numerical differences between the reference deployment and any other deployment are simple  
273 to compute with any variety of binary difference utility, such as the Unix `cmp` or `checksum`, or the  
274 NetCDF differencing operator `ncdiff`. In all testing so far and regardless of computing system, all  
275 of the WRF output files from the same configuration are bit-wise identical.

276 A container that builds the WRF internally, via the gnu compilers that are inside the container,  
277 can equally be deployed on any platform that supports containers. Experiments have shown that  
278 those results are also identical to the output from pre-built binary executables deployed as above.

279 Container deployment on individual servers, regardless of the number of CPUs or computational  
280 cores, also yield identical results. Both the container technology, and also that the WRF code is  
281 sufficiently mature to make the same calculations on any parallel topology, are responsible for  
282 that result. Numerical reproducibility across different parallel topologies is non-trivial for such a  
283 complex code (e.g. Thomas et al. 2002), and the WRF developers deserve credit. We have not  
284 tested all possible configurations of the WRF, which number in the tens of thousands, so cannot  
285 guarantee that all configurations are robust to differences in parallel topologies.

286 Results from the cross-platform testing here show that numerical reproducibility is easy to  
287 achieve with containers, within some easily controlled limitations that are discussed next. By  
288 wrapping all dependencies in a container, researchers working on disparate local hardware, OS,  
289 and compiler resources can obtain numerically identical results.

290 Containers are an attractive upgrade from virtual machines because the computational overhead  
291 is significantly less. In an extensive analysis, Felter et al. (2014) found that configured properly,  
292 Docker containers surpassed Kernel-based Virtual Machines (KVM) in every performance test.

293 Computational and memory performance overhead is very small, and any performance penalty is  
294 on I/O and interactions with the operating system. Although OS interactions are negligible for  
295 NWP models, I/O can be a limitation. Fortunately the overhead is dependent on the number of I/O  
296 operations, not the I/O size. In many relevant problems the WRF does not output many files, but  
297 it can output large files for large computational domains.

298 In tests on the Sandy simulation we also find minimal performance impact from running within  
299 the Docker container. Identical binary executables were run four times on the same machine both  
300 inside the container, and outside the container on a Linux server. Results showed an average 3.6%  
301 overhead on wall-clock time when running in the container. The overhead will be greater for  
302 operating systems requiring a VM layer (e.g. Windows) with older versions of Docker.

303 This Sandy simulation requires relatively few floating-point operations, compared to model ex-  
304 ecutions in typical research and operational NWP today. It is parallelized on shared-memory  
305 computer servers. On a Macbook Air as described in the second line of Table 1, the sequence of  
306 pulling the container from the repository, running the WRF, and making plots takes approximately  
307 4 minutes. The example demonstrates the potential for classroom utility.

308 Performance reproducibility can be contrasted with the numerical reproducibility examined  
309 above. Performance reproducibility deals with the issue of obtaining the same performance (run  
310 time, throughput, latency, etc.) across executions. Reproducing performance is important for gain-  
311 ing accurate evaluations of that performance. For example, when performance improvements to  
312 the WRF code base are implemented, evaluating improvements is difficult because it is difficult  
313 to account for the changes in hardware over time. In Jimenez et al. (2016), the authors introduce  
314 a framework to control the cross-platform variability arising from changes in CPU architectures  
315 and generations. This technique leverages Linux’s cgroups technology (an underlying component  
316 of Docker), which allows a user to specify the CPU bandwidth associated with a container. For

example, one can constrain a WRF container to have 50% the capacity of a CPU core (instead of 100% of the bandwidth available, the default behavior). With this feature, one can find the amount of CPU capacity that emulates the behavior of a slower (usually older) machine. Table 2 shows the results of applying this technique to reproduce the performance of the Hurricane Sandy simulation on multiple platforms.

In Table 2, machine `issdm-6` is the base system where a reference simulation was executed. When this same simulation is re-executed in other (newer) machines, an inherent speed-up in run-time results from the improvements in newer CPU technology. For example, machine `dwill` improves the run time of the original execution by approximately 2.8 times. Using the CPU calibration techniques introduced in Jimenez et al. (2016), which leverage the OS-level virtualization in containers, we carefully constrain the CPU of the machines and re-execute the simulation. As described in Jimenez et al. (2016), the limits on CPU are obtained by executing a battery of micro-benchmarks that characterize the CPU performance of a base system (`issdm-6` in this case). Benchmark results are used to calibrate the target machine (for example `dwill`) by finding the percentage of CPU capacity it needs to execute the same list of micro-benchmarks and emulate the performance of the base system. In Table 2, the `CPU limits` column denotes the tuned amount of CPU capacity that the machine was given for a simulation, with the goal of emulating the performance behavior of the base machine. For the case of `dwill`, the limitations on CPU bandwidth (38% of CPU capacity) bring the run-time down (introduce an artificial slowdown), closely resembling the original performance.

Performance reproducibility also enables planning for computing needs. This is becoming more important as cloud computing becomes more prevalent. Commercial cloud vendors such as Amazon Web Services, Packet, Google Kubernetes ([kubernetes.io](https://kubernetes.io)), and Microsoft Azure ([azure.microsoft.com](https://azure.microsoft.com)) generally support container deployment. A researcher or company writ-

341 ing a proposal that may include a need for cloud computing resources can get accurate estimates  
342 of performance, and expect it to hold on the variety of hardware offered by vendors.

343 This section presented a subset of results from runs on many different machines. Although parallel  
344 WRF runs have been part of the testing, they are limited so far to single shared-memory nodes.  
345 Distributed (multiple-nodes or servers) Message Passing Interface (MPI) for WRF within Docker  
346 containers is under investigation. Progress depends on the MPI-based WRF application having  
347 awareness of other linked WRF executables in other containers that may be part of a collective  
348 container group. Currently MPI does not have that capability, although Docker containers can be  
349 made aware of parallel containers within a collective.

## 350 **5. Uncertainty quantification**

351 A corollary to numerical reproducibility is that a precise quantification of uncertainty from other  
352 factors is enabled. The reproducibility means that by leveraging available computer power whenever  
353 ever it is available, hundreds or thousands of simulations can be produced with perturbations  
354 introduced in initial conditions or model configurations. Such an ensemble would reflect only the  
355 perturbations given to it, and not be contaminated by random or systematic errors introduced from  
356 traditional computing platforms or changing compiler options. Leaving ensemble investigations  
357 for future work, here we examine the effects of compiler optimization. The effects we report here  
358 are certain to be reproducible across many computing platforms.

359 A simple pair of simulations illustrate the effects of compiler optimization on numerical pre-  
360 dictions. The basic FORTRAN optimization flag for the gnu compiler was successively set to  
361 levels from 0 to 3, which is a typical range for compilers. Setting -O0 forces the compiler to avoid  
362 optimizing calculations and memory access, and -O3 allows the compiler a number of relatively  
363 aggressive optimization strategies. Those may change the order of operations, for example. Dif-

364    ferent compilers introduce different optimizations at that level. We found in this specific case that  
365    levels 0 to 2 led to a WRF executable that produced identical results, but level three produces  
366    differences. One of the optimizations introduced at level 3 has an effect on output, but it is beyond  
367    the scope of this work to determine exactly what that is. Instead, we simply compare the output  
368    between levels 0 and 3.

369    A histogram of 3D grid-point differences between the two Sandy simulations, after 12 h of  
370    simulation time, shows that meaningful differences are present (Fig. 3). Although the majority  
371    of differences are small, the distribution tails indicate the onset of local perturbations. We know  
372    those perturbations will grow and propagate up-scale because the model equations are chaotic (e.g.  
373    Lorenz 1969). We can expect that higher-resolution grids, subject to more small-scale nonlinear  
374    processes, will lead to greater rates of growth of the differences between the two simulations.  
375    Grids covering larger regions will be subject to greater rates of growth too, because the dynamical  
376    interactions cover a broader range of scales.

## 377    **6. Discussion**

378    This article presents a set of open-source Docker containers intended to provide a basis for a  
379    WRF execution, and ancillary workflow components. As shown above, the container infrastruc-  
380    ture offers many advantages for enabling research and education to be reproducible and more  
381    collaborative. A few more discussion points are summarized here.

382    The containers offer numerical and performance reproducibility in an easily, and rapidly, deploy-  
383    able framework that can enable collaborative research and education. In section 5 we demonstrated  
384    the effects of compiler flags. Many NWP experiments reported in the literature focus on small dif-  
385    ferences between simulations, which are interpreted as physically relevant signals to be diagnosed.  
386    Although the majority of the differences may be physically relevant, the reproducibility provided

387 by containers offers certainty. The extension is that uncertainty from sources besides numerical  
388 truncation can be isolated. To further understand these issues, a study evaluating whether com-  
389 piler optimizations, hardware variations, and parallel topologies lead to random-like or systematic  
390 errors would be helpful.

391 The containers enable the same kind of collaborative environment as a single community re-  
392 source, but by making use of any resources locally available to individual researchers. In general,  
393 identical compilers and operating systems are not available to collaborating researchers, except  
394 when all participants are working on the same community computing resource and in the same  
395 environment. Numerical reproducibility even within collaborating groups is not straightforward.  
396 With containers, collaboration among groups of scientists that need to work on the same code, or  
397 produce parallel simulations with only small variations, is immediately enabled. Examples include  
398 ensemble forecasts run on distributed and inhomogeneous platforms by multiple people in differ-  
399 ent locations, model developers from different institutions who are working to improve a physics  
400 scheme, or researchers collaborating to better quantify an energy budget within the model.

401 Educational activities clearly benefit from containers. Hands-on exercises are a valuable part of  
402 any student learning experience. Containers give instructors control over what parts of the mod-  
403 eling system students will experience directly. For example the WRF could be run initially as a  
404 “black box,” and then intermediate steps made accessible as students work through exercises. Soft-  
405 ware containers help standardize modules that can be easily shared amongst teaching colleagues  
406 with vastly different technology environments.

407 The broader science community has recently been placing a greater emphasis on reproducible  
408 research, and containers offer a key step toward reproducible NWP research. Published papers can  
409 cite a specific container version used in one part of a research flow (e.g. the model, set of analysis  
410 tools, or a data container with initialization files). As long as that container revision exists and is

411 accessible, research consumers can reproduce the results. This provides not only an unprecedented  
412 level of openness in NWP research, but also an easier way for researchers to build on published  
413 results. The resources below in section 7 offer one component of the suite of tools needed to  
414 enable fully reproducible science. Metadata and analysis methods also need to be tracked and  
415 made available. Tools for that exist, and should be adopted by the NWP community as we go  
416 forward.

## 417 7. Resources

418 This section provides links to on-line content, including documentation and open repositories for  
419 WRF-based containers. A basic knowledge of repositories, and access to GitHub and DockerHub  
420 public repositories, are needed to pull source code or container images from the repositories. A  
421 list of resources follows:

422 Docker-WRF project home page: [www.ral.ucar.edu/projects/ncar-docker-wrf](http://www.ral.ucar.edu/projects/ncar-docker-wrf)

423 Slack channel (user discussion forum): [ncar-dockerwrf.slack.com](https://ncar-dockerwrf.slack.com)

424 WRF/WPS container image: [hub.docker.com/r/bigwxwrf/ncar-wrf/](https://hub.docker.com/r/bigwxwrf/ncar-wrf/)

425 Static geography data container image: [hub.docker.com/r/bigwxwrf/ncar-wpsgeog/](https://hub.docker.com/r/bigwxwrf/ncar-wpsgeog/)

426 Input data container for Sandy simulation: [hub.docker.com/r/bigwxwrf/ncar-wrfinputsandy/](https://hub.docker.com/r/bigwxwrf/ncar-wrfinputsandy/)

427 Input data container for Katrina simulation: [hub.docker.com/r/bigwxwrf/ncar-wrfinputkatrina/](https://hub.docker.com/r/bigwxwrf/ncar-wrfinputkatrina/)

428 NCL script container for producing images: [hub.docker.com/r/bigwxwrf/ncar-ncl/](https://hub.docker.com/r/bigwxwrf/ncar-ncl/)

429 GitHub repository Docker files and scripts to build images: [github.com/ncar/docker-wrf](https://github.com/ncar/docker-wrf)

430 *Acknowledgments.* Partial funding for this work is from the following: National Science Foun-  
431 dation awards ATM0753581/M0856145 to NCAR, and 1450488 to the University of California  
432 Santa Cruz. Sandia National Labs, and LANL/UCSC Institute for Scalable Scientific Data Man-  
433 agement (ISSDM) also contributed funding. Work at the University of North Dakota was funded

434 by NSF ACI-1450168 and the North Dakota Space Grant Consortium. Cindy Halley-Gotway at  
435 NCAR produced the schematic in Fig. 1. Amazon Web Services is acknowledged for providing  
436 an educational resource grant used for some of these investigations.

437 **References**

- 438 Baker, A. H., and Coauthors, 2015: A new ensemble-based consistency test for the Community  
439 Earth System Model (pyCECTv1.0). *Geosci. Model Devel.*, **8**, 2829–2840, doi:doi:10.5194/  
440 gmd-8-2829-2015.
- 441 Felter, W., A. Ferreira, R. Rajamony, and J. Rubio, 2014: An up-  
442 dated performance comparison of virtual machines and Linux contain-  
443 ers. Tech. Rep. RC25482, IBM Research Report. Currently available from  
444 [http://domino.research.ibm.com/library/cyberdig.nsf/papers/0929052195DD819C85257D2300681E7B/\\$](http://domino.research.ibm.com/library/cyberdig.nsf/papers/0929052195DD819C85257D2300681E7B/$)
- 445 Jimenez, I., C. Maltzahn, J. Lofstead, A. Moody, K. Mohror, R. Arpaci-Dusseau, and A. Arpaci-  
446 Dusseau, 2016: Characterizing and reducing Cross-Platform performance variability using OS-  
447 level virtualization. *The First IEEE International Workshop on Variability in Parallel and Dis-*  
448 *tributed Systems (VarSys 2016) (VarSys 2016)*, Chicago, USA.
- 449 Lorenz, E. N., 1963: Deterministic nonperiodic flow. *J. Atmos. Sci.*, **20**, 130–141.
- 450 Lorenz, E. N., 1969: The predictability of a flow which possesses many scales of motion. *Tellus*,  
451 **21**, 289–307.
- 452 Skamarock, W. C., and Coauthors, 2008: A description of the advanced research WRF Version 3.  
453 Tech. Rep. TN-475, National Center for Atmospheric Research.
- 454 Thomas, S., J. P. Hacker, M. Desgagné, and R. B. Stull, 2002: An ensemble analysis of forecast  
455 errors related to floating point performance. *Wea. and Forecasting*, **17**, 898–906.

456 **LIST OF TABLES**

457	<b>Table 1.</b> Representative sample of container test platforms for the Hurricane Sandy sim-	
458	ulations. An asterisk denotes a deployment on an Amazon EC2 resource; the	
459	others are on a laptop or desktop/server. A double asterisk denotes a deploy-	
460	ment on a Packet Tiny Atom resource. . . . .	23
461	<b>Table 2.</b> Run time of container executions for the Hurricane Sandy simulations on mul-	
462	tiple machines. The normalized run time is the ratio between the reference	
463	system <code>issdm-6</code> and the platform identified in the particular row. 100% de-	
464	notes no limitations imposed on the CPU, i.e. the simulation had all the CPU	
465	capacity from the system. . . . .	24

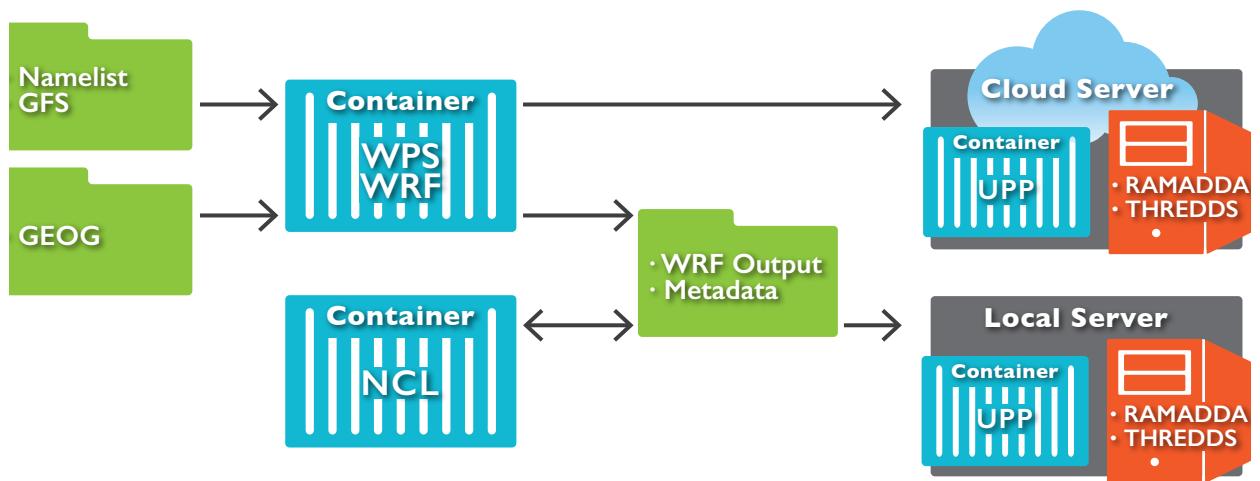
466 TABLE 1. Representative sample of container test platforms for the Hurricane Sandy simulations. An asterisk  
 467 denotes a deployment on an Amazon EC2 resource; the others are on a laptop or desktop/server. A double  
 468 asterisk denotes a deployment on a Packet Tiny Atom resource.

OS	Chip/CPU	Cores
OS X 10.10.5	Intel Xeon E5	6
OS X 10.9.5	Intel Core i7	2
Ubuntu 14.04	AMD Opteron 6320	16
Ubuntu 14.04	Intel Xeon E5	16
RHEL	Intel Xeon X5550	8
*Ubuntu 14.04	Intel Xeon E5-2666 v3	32
**Ubuntu 14.04	Intel Atom C2550	4

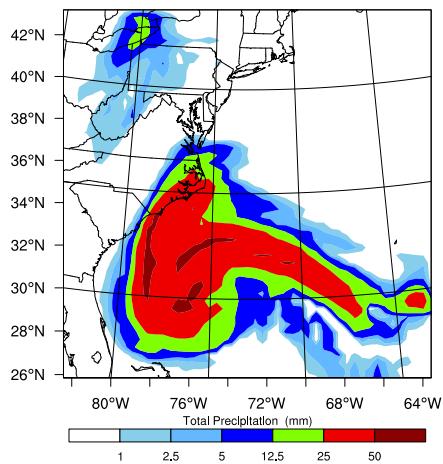
469 TABLE 2. Run time of container executions for the Hurricane Sandy simulations on multiple machines.  
 470 The normalized run time is the ratio between the reference system `issdm-6` and the platform identified in the  
 471 particular row. 100% denotes no limitations imposed on the CPU, i.e. the simulation had all the CPU capacity  
 472 from the system.

Machine ID	CPU Model	Runtime	CPU limit	Normalized
issdm-6	AMD Opteron 2212	1251.697012	100%	1.000000
packet0	Intel Atom C2550	1173.848259	100%	1.066319
packet0		1265.819255	99%	0.988843
nibbler	Intel i7 930	571.163721	100%	2.191485
nibbler		1425.671584	63%	0.877970
dwill	Intel i5 2400	442.803647	100%	2.826754
dwill		1244.367972	38%	1.005890
rackform4	AMD Opteron 6320	689.738835	100%	1.814741
rackform4		1560.987534	62%	0.801862
pl2	Intel E5-2630 V2	425.031332	100%	2.944952
pl2		1431.383836	48%	0.874466
node	Intel E5-2630 V3	396.553413	100%	3.156440
node		1564.413064	56%	0.800106

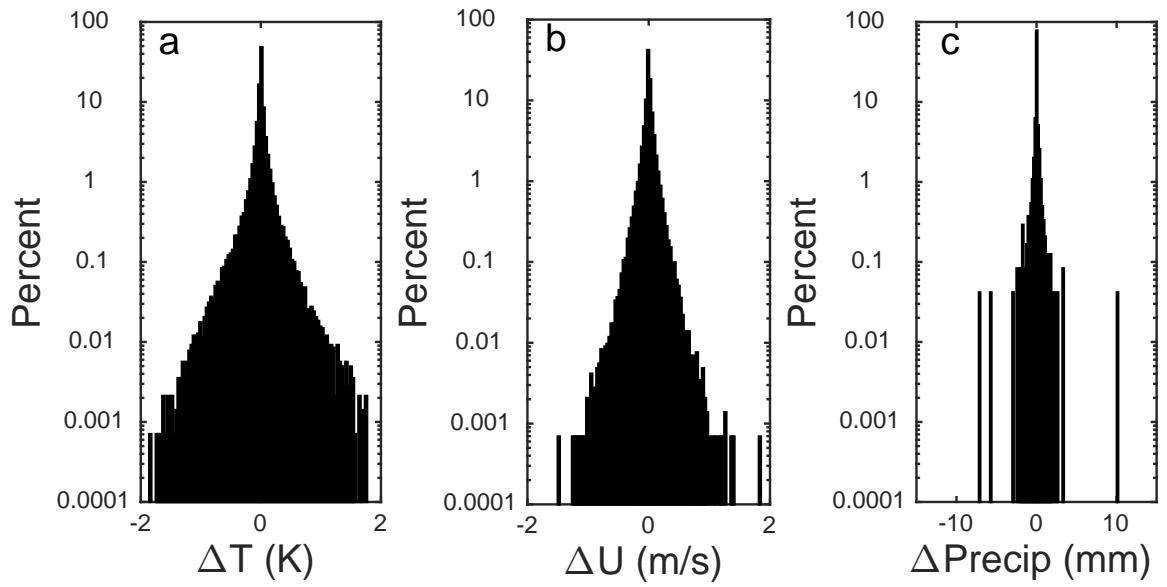
473 LIST OF FIGURES



488 FIG. 1. Schematic illustrating an example deployment of the WPS and WRF in a Docker container. Green  
 489 folders denote files on a local file system and which are easily made available to the containers or a local user.  
 490 This example uses Global Forecast System (GFS) output files for WRF initial conditions and lateral boundary  
 491 conditions. The containers (blue) interact with the local file system along the paths shown with arrows. Output  
 492 and metadata can be stored and served with local resources, or in a hosted computing and storage environment  
 493 (the cloud). Data can be served in either case with help from, for example, the Geode Systems RAMADDA or  
 494 Unidata Thematic Real-time Environmental Distributed Data Services (THREDDDS) software.



495 FIG. 2. Simulated 12-h total precipitation from the demonstration Hurricane Sandy simulation, valid 0000  
496 UTC 28 Oct 2012.



497 FIG. 3. Differences between grid-point values of two Sandy simulations, (a) temperature and (b) zonal wind,  
 498 valid 0000 UTC 28 Oct 2012 after 12 h. WRF executables differ only in a single compiler optimization flag. All  
 499 grid points in the 3D volume are included.